

BIOS 635: Bias and variance trade-off, Classification, K-nearest neighbor

Kevin Donovan

1/26/2021

Review

- Homework 1 assigned, due on 1/28 at 11PM through GitHub Classroom
- Article Evaluation 1 assigned, due on 2/2 through GitHub Classroom
- Office Hours: Wednesday 10-11AM
- Last week: discussed supervised and unsupervised learning, curse of dimensionality, evaluating model performance

Supervised Learning

- Features:

$$X = \begin{pmatrix} X_1 \\ X_2 \\ \dots \\ X_p \end{pmatrix}$$

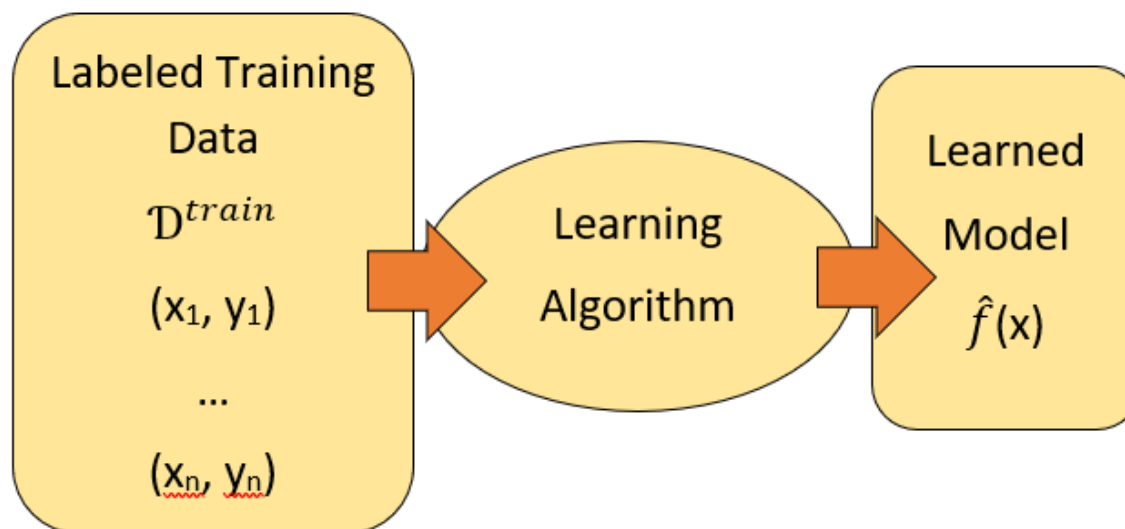
- Model:

$$Y = f(X) + \epsilon$$

- Goals

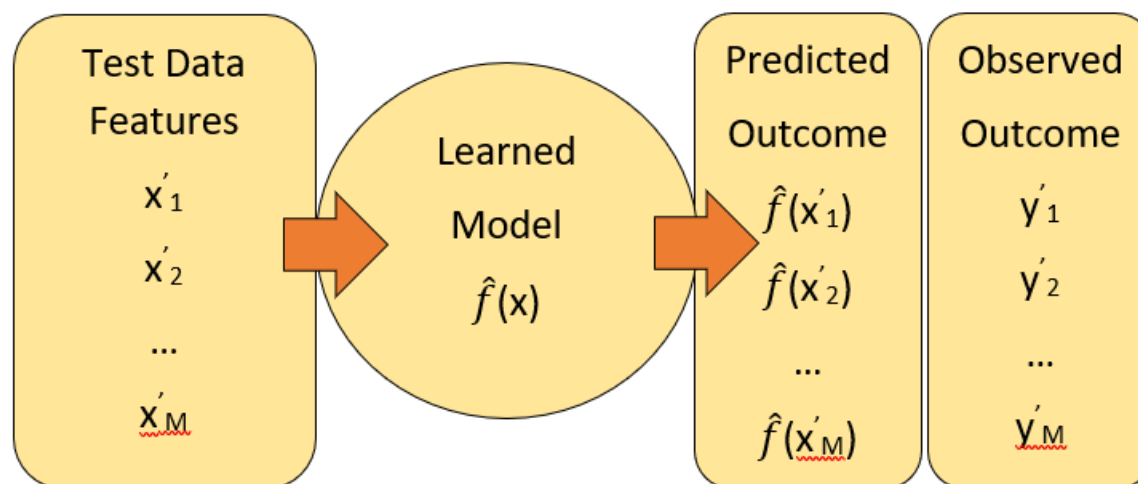
- Define $f(X)$; e.g. $f(X) = E(Y|X = x)$
- Model and estimate $f(X)$, denoted $\hat{f}(x)$
- Define metric to evaluate estimated model; e.g. $MSE(x) = E[(Y - \hat{f}(X))^2|X = x]$

Supervised Learning



- Give learner training data
- Learner returns model $\hat{f}(x)$

Supervised Learning



- Give test data estimated model $\hat{f}(x)$
- Compare predicted outcomes from $\hat{f}(x)$ with observed

Mean Squared Error Decomposition

Recall: MSE for estimate at $X = x$ can be decomposed into

$$MSE_{\hat{f}}(x) = E[(Y - \hat{f}(X))^2 | X = x] = [f(x) - \hat{f}(x)]^2 + Var(\epsilon)$$

Consider taking expectation marginally (i.e., across Y and X).

Can show

$$E[(Y - \hat{f}(X))^2] = E_x[\text{bias}(\hat{f}(x))^2] + E_x[\text{Var}(\hat{f}(x))] + \text{Var}(\epsilon)$$


where $\text{bias}(\hat{f}(x)) = E[\hat{f}(x)] - f(x)$

Bias-variance trade-off


Above means bias **and** variance of model increases expected model error

Creates tradeoff:

- Higher complexity: **decreased** bias but **increased** variance
- Lower complexity: **increased** bias but **decreased** variance



Bias-variance trade-off




Classification

Suppose instead response Y is categorical

e.g. cancer stage is one of $C = (0, 1, 2, 3, 4)$ where 0 indicates cancer-free

Goals:

- Build classifier $\hat{f}(X)$ that maps a category from C to future observation X
- Assess uncertainty in each classification
- Understand roles of predictions $X = (X_1, X_2, \dots, X_p)$



Classification

What to model?:

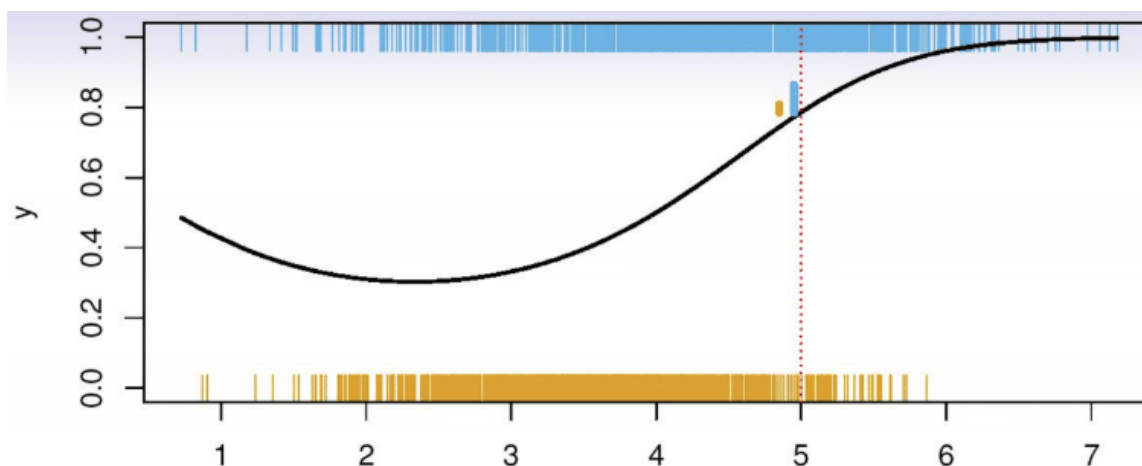
Let $p_k(x) = \Pr(Y = k|X = x)$, $k = 1, 2, \dots, K$

Denoted as the **conditional class probabilities** at x

If these are known, can define classifier at x by

$$f(x) = j \text{ if } p_j(x) = \max[p_1(x), \dots, p_K(x)]$$

Denoted as the **Bayes optimal classifier** at x



Classification metrics

Basic:

$$\text{accuracy} = \frac{\# \text{ correct predictions}}{\# \text{ test instances}}$$

$$\text{error} = 1 - \text{accuracy}$$

These are in general **not sufficient** (why?)

Confusion Matrix

		Predicted Class	
		Yes	No
Actual Class	Yes	TP	FN
	No	FP	TN

$$\text{accuracy} = \frac{TP + TN}{P + N}$$

- Want to correctly identify positive **and** negative instances accurately
- For positive instances, have the following metrics:

$$PPV = \frac{TP}{TP + FP}$$


$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2}{\frac{1}{PPV} + \frac{1}{\text{Sensitivity}}}$$

Confusion Matrix


Example

During the COVID-19 pandemic, different metrics to quantify risk:



COVID-19 Deaths Per 100,000 Inhabitants: A Comparison

COVID-19 deaths per 100,000 of the population in the 10 worst affected countries*



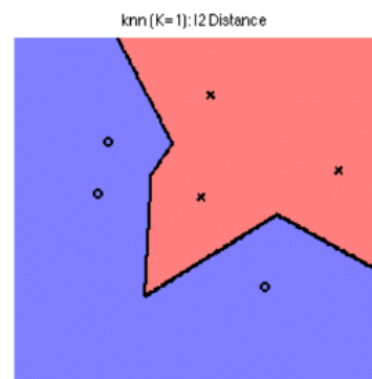
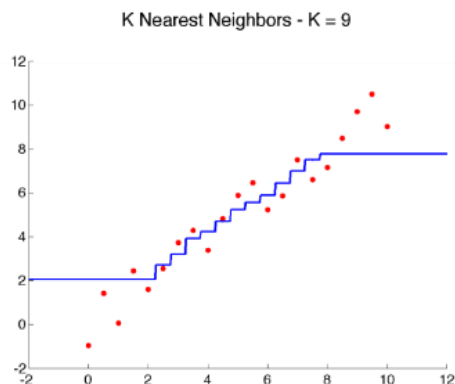
* As of August 09, 2020 at 03:00 AM EDT
Source: Johns Hopkins University



K-nearest Neighbor (KNN)


Simple and **flexible** algorithm:

1. Given training data $D = \{\mathbf{x}_i, y_i\}$, distance function $d(\cdot, \cdot)$ and input \mathbf{x}
2. Find $\{j_1, \dots, j_K\}$ closest examples with respect to $d(\mathbf{x}, \cdot)$
 - (regression) if $y \in \mathbf{R}$, return average: $\frac{1}{K} \sum_{k=1}^K y_{j_k}$
 - (classification) if $y \in \pm 1$, return majority: $\text{sign}(\sum_{k=1}^K y_{j_k})$




Distance Metrics

- Different metrics can dramatically change neighborhoods:



$$\text{Dist}(\mathbf{a}, \mathbf{b}) = (a_1 - b_1)^2 + (a_2 - b_2)^2$$



$$\text{Dist}(\mathbf{a}, \mathbf{b}) = (a_1 - b_1)^2 + (3a_2 - 3b_2)^2$$

- Standard Euclidean distance metric:

- 2D: $\text{Dist}(a, b) = \sqrt{[(a_1 - b_1)^2 + (a_2 - b_2)^2]}$
- Multidim: $\text{Dist}(a, b) = \sqrt{\sum (a_i - b_i)^2}$

Distance Metrics


$$\text{Dist}(a, b) = \left(\sum_i |a_i - b_i|^p \right)^{1/p}$$

$p = 1$, Manhattan Distance


$p = 2$, Euclidean Distance

$p = \infty$, Chebychev Distance


Euclidean Distance



Chebyshev Distance



Manhattan Distance



$$\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2} \quad \max(|x_1 - x_2|, |y_1 - y_2|) \quad |x_1 - x_2| + |y_1 - y_2|$$

KNN Examples

I. Regression

- Suppose we want to predict the number of wins for a NBA team

```
## k-Nearest Neighbors
##
## 30 samples
## 2 predictor
##
## Pre-processing: centered (2), scaled (2)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 30, 30, 30, 30, 30, ...
## Resampling results across tuning parameters:
##
## k    RMSE      Rsquared    MAE
## 5    4.585041  0.8918501  3.700367
## 7    4.845494  0.8979390  3.972826
## 9    5.044867  0.8978793  4.163350
## 11   5.408245  0.9069323  4.400432
## 13   6.119488  0.8964707  4.995359
## 15   6.908745  0.8930579  5.724303
## 17   7.706909  0.8906881  6.415243
## 19   8.613891  0.8744222  7.197625
## 21   9.406709  0.8592141  7.902947
## 23  10.066420  0.8578698  8.499900
## 25  10.842491  0.7900220  9.237286
## 27  11.829461  0.7020191 10.111490
## 29  12.466729  0.6103163 10.753089
## 31  12.744056      NaN 11.034206
## 33  12.744056      NaN 11.034206
## 35  12.744056      NaN 11.034206
## 37  12.744056      NaN 11.034206
## 39  12.744056      NaN 11.034206
```

```
## 41 12.744056      NaN 11.034206
## 43 12.744056      NaN 11.034206
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 5.
```


KNN Examples

2. Classification

- Suppose we want to predict whether a NBA team made the playoffs


```
## k-Nearest Neighbors
##
## 30 samples
## 2 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (2), scaled (2)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 30, 30, 30, 30, 30, 30, ...
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##  5  0.8625986  0.7151352
##  7  0.8640695  0.7320667
##  9  0.8463797  0.6966267
## 11  0.8348877  0.6892596
## 13  0.8283119  0.6807602
## 15  0.8381627  0.7011166
## 17  0.7941796  0.6165274
## 19  0.7428680  0.5514878
## 21  0.6846963  0.4460085
## 23  0.6910537  0.4625074
## 25  0.5882703  0.2921750
## 27  0.4964599  0.1476145
## 29  0.4181164  0.0000000
## 31  0.4181164  0.0000000
## 33  0.4181164  0.0000000
## 35  0.4181164  0.0000000
## 37  0.4181164  0.0000000
```

```
## 39 0.4181164 0.0000000
## 41 0.4181164 0.0000000
## 43 0.4181164 0.0000000
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 7.
```




KNN Examples

KNN: $K=1$



KNN: $K=100$



Train and Test Error

